

Intro. To Deep Learning

Phuriwat Angkoondittaphong

What is deep learning

- Deep learning is subset of machine learning that use Artificial Neural Network (ANN).
- Machine learning is an AI algorithm that can learn from data using statistic, or something else.

How (mathematical) model learn


- Data
- Example model
- Objective function
- Optimization
 - Analytically
 - Numerically

What is Data

- Consider regression problem

$$D = \{(x_1, y_1), (x_2, y_2), \dots, \}$$

where $x_i, y_i \in \mathbb{R}$



Consider a linear regression model

$$M : \mathbb{R} \rightarrow \mathbb{R}$$
$$M(x) = wx + b$$

- The model has 2 parameters, w and b
- By adjusting, it means to change parameters so that it fit to the training data more
- M is estimator, meaning, the model tried to estimate some value on the real line.


How to indicate “best” model (Objective function)

- how do we know that the model perform good or bad
- One idea is to sum up the error
- One of them called sum of squared error

$$e = \sum_i (y_i - M(x_i))^2$$


Then we can find m and b such that produce minimum e

$$(\hat{m}, \hat{b}) = \operatorname{argmin}_{(m,b)} e$$



Optimizing Linear Regression Analytically

- Since, the error function is quite simple, we can solve for a minima in the function using calculus of variations. Leave this as exercise for viewer.



At $e' = 0$, the e will be its minimum

Optimizing Linear regression: Deep learning perspective

- But in deep learning, we usually use model that is more complex than this
- Something like 1M of parameters, which can be hard to calculate derivative analytically.
- However, we can calculate it numerically.
- The algorithm is called Gradient descent

Graph of e over a variable

If e' at this point, we know that if we increase variable, e will be reduce

$$e' < 0$$

If e' at this point, we know that reduce variable to reduce e

$$e' > 0$$

$$e' = 0$$

Optimal point

Variable


With this, the variable knows its direction that can make the model better

Learning rate

If e' at this point, increase the variable to reduce e .


We don't know how much we should take a step.

Example of a step



Learning rate

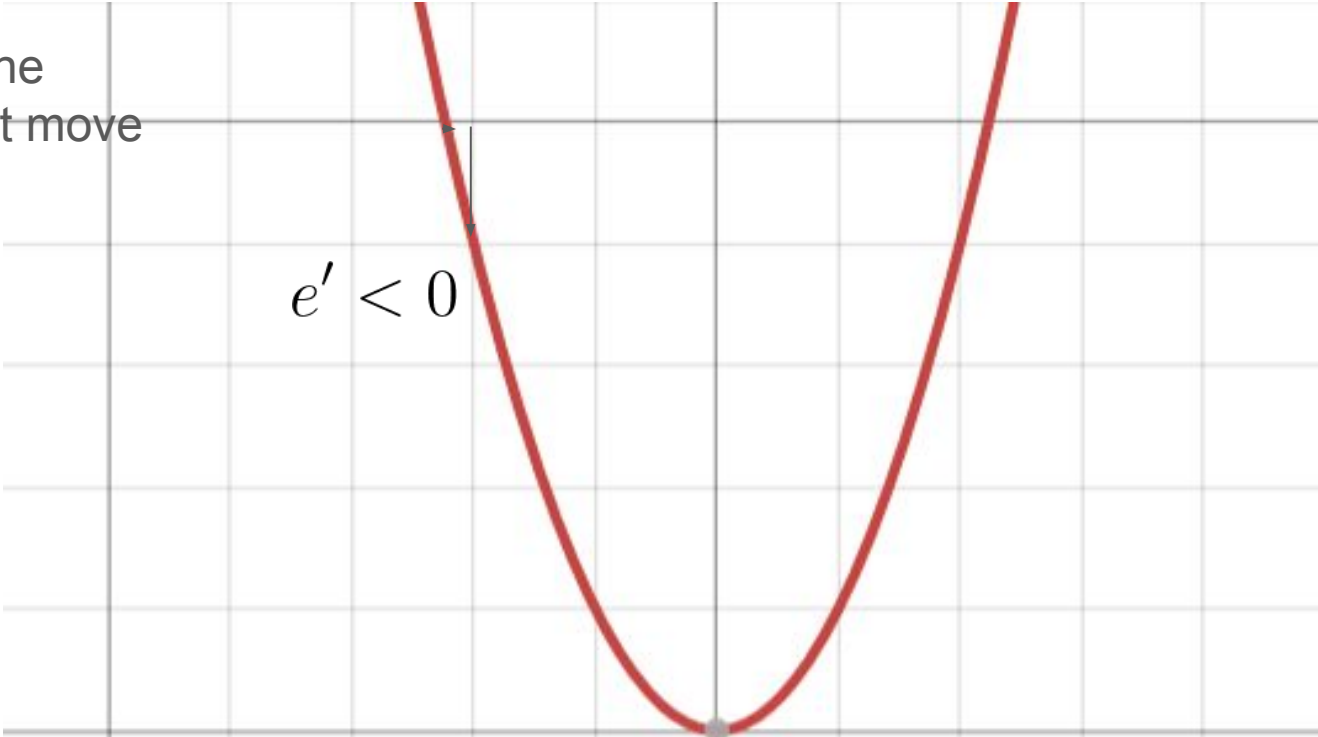
The step should be
not too large, not too small



Learning rate

Example of small step

Too small and the
parameter won't move



$e' < 0$

Learning rate


- Size of the step is usually called Learning rate (LR)
- It is one of the most important hyperparameter

How we calculate derivative (gradient)

- It use computational graph something, something, something, ...
- A story for another day.
- Today we gonna use PyTorch to calculate for us

PyTorch

- PyTorch is a library for scientific computing.
- Written in C++ and CUDA with Python interface
- It provides fast multidimensional array computing functions (eg. matrix multiplication) on both CPU and GPU.
- It also provides easy way to compute gradient (automatic differentiation system).



PyTorch installation

- Installation (not recommended)

PyTorch Build	Stable (2.2.2)		Preview (Nightly)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 11.8	CUDA 12.1	ROCm 5.7	CPU
Run this Command:	<code>pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118</code>			

- If you dont want to install it yourself, just use google colab

PyTorch Basic Workshop

Train simple linear regression

	A	B	
1	Quiz	Final	
2	25.2	30.5	
3	8.4	13.6	
4	4.9	14	
5	11.1	16.6	
6	9	15.8	
7	7.8	6.3	
8	3	2.4	
9	16	27.7	
10	2.5	0	
11	13.8	16.8	
12	9.6	14.9	
13	11.8	7.6	
14	16.9	25.9	

Model training in PyTorch

1. Create dataset
2. Define the model
3. Choose optimizer
4. Choose loss function
5. Train
6. Evaluate

Create dataset

```
1 class ScoreDataset(torch.utils.data.Dataset):
2     def __init__(self, x, y):
3         self.x = x
4         self.y = y
5         assert len(x) == len(y)
6         self.n = len(x)
7     def __getitem__(self, idx):
8         return self.x[idx], self.y[idx]
9     def __len__(self):
10        return self.n
```


Define model

- In order to compute gradient, the model must consist of differentiable components.


Model Layers: Linear Transformation

- Linear
 - The simplest layer of all nn.
 - “Linear” from Linear transformation
 - Simply Linear regression on steroid

$$y = Wx + B$$





Linear regression with one input and one output



Linear regression with 3 inputs and 2 outputs

Model Layers: Activation function

- Chaining many linear layers do not work.
- We need non-linearity between linear layers.
- The most common one is ReLU = $\max(0, x)$
- Sigmoid $\sigma(x) = \frac{1}{1 + e^{-x}}$
- SoftMax (usually used on the last layer of classification)



Training

```
def train_model_one_epoch(model, train_dataloader, optimizer, loss_fn):
    cumulative_loss = 0
    for data in train_dataloader:
        # Every data instance is an input + label pair
        inputs, labels = data

        # Zero your gradients for every batch!
        optimizer.zero_grad()

        # Make predictions for this batch
        outputs = model(inputs)

        # Compute the loss and its gradients
        loss = loss_fn(outputs, labels)
        loss.backward()

        # Adjust learning weights
        optimizer.step()

        cumulative_loss += loss.item()

    cumulative_loss /= len(train_dataloader)
    return cumulative_loss
```


Train simple classifier

